

1 Problem 2 Take Home Final Exam

```
# Importing the necessary libraries
import matplotlib.pyplot as plt
import numpy as np
from numpy.linalg import norm

np.random.seed(999) # Seed random number generator for consistency

#Initialize the matrix dimensions
n = 100
m = 200

# Create a random problem
A = np.random.randn(m,n)

# Pick the stopping condition
eta = 1e-4
eps = 1e-8
i_max = 1000
```

1.1 Part a - Gradient Descent

```
# Function goes here
def f(x):
    return -sum(np.log(1-np.dot(A,x))) - sum(np.log(1-x)) - sum(np.log(1+x))

# Have to differentiate and hard code derivative
def grad_f(x):
    s1 = 1/(1-np.dot(A,x))
    s2 = np.dot(np.transpose(A),s1)
    return s2+1/(1-x)-1/(1+x)
```

1.1.1 Algorithm for Backtracking Line Search

while $f(x + t \text{ delta } x) > f(x) + \alpha t \text{ gradient}(f(x)) \text{ delta } x$: $t = \beta t$

```
# Backtracking method for updating t
def backtrack(x,dx,alpha,beta):
    t=1.0
    # Make sure updating to feasible x, use approximation to ensure in feasible domain
    s = x+np.dot(t,dx)
    while (np.amax(np.dot(A,s))>1) or (np.amax(abs(s))>1):
        t=beta*t
        s = x+t*dx
    # Backtrack in feasible domain
    s2 = np.dot(np.transpose(grad_f(x)),dx)
    while f(s) > f(x)+alpha*t*s2:
        t=beta*t
        s = x+t*dx
        s2 = np.dot(np.transpose(grad_f(x)),dx)
    return t
```

1.1.2 Algorithm for gradient descent

1. $Dx = -1 * \text{gradient } f(x)$
2. Line search - choose step size t via backtracking method
3. Update $x = x + t Dx$

```
def gradient_descent(iterations, alpha, beta):
    y = np.array([])
    s = np.array([])
    # Use zero as initial guess
    x = np.zeros(n)
    x = x.reshape(n,1) # Make sure x is a vector
    # Repeat
    for i in range(iterations):
        # Step 1
        y = np.append(y, f(x))
        dx = -1*grad_f(x)

        # Step 2
        t = backtrack(x, -grad_f(x), alpha, beta)
        s = np.append(s, t)

        # Step 3
        x = x + t*dx

        p = f(x)

        if norm(grad_f(x)) < eta:
            break

    return y, s, p
```

Making a plot function to play with alpha and beta more easily

```
def gradient_descent_plot(alpha, beta):
    # Run gradient descent algorithm
    y, s, p = gradient_descent(i_max, alpha, beta)

    fig, ax = plt.subplots(2, 1)

    ax[0].semilogy(y-p)
    ax[0].set_title('Convergence to Optimal Solution')
    ax[0].set_xlabel('Iterations')
    ax[0].set_ylabel('f - p_star')

    ax[1].stem(s)
    ax[1].set_title('Step Length')
    ax[1].set_xlabel('Iterations')
    ax[1].set_ylabel('t')

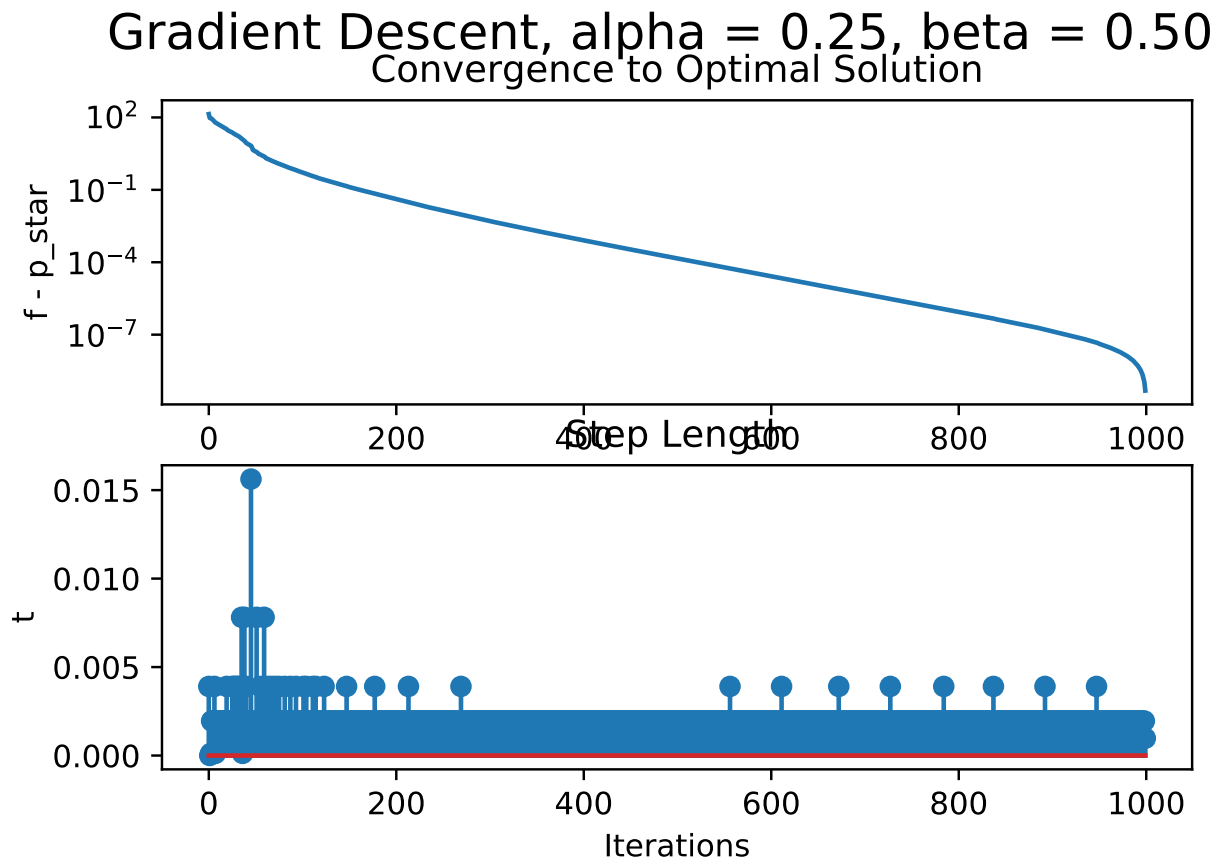
    fig.suptitle('Gradient Descent, alpha = %1.2f, beta = %1.2f'%(alpha, beta), fontsize=12)

    return fig, ax
```

1.1.3 Figures and analysis

Below is the code to perform gradient descent with $\alpha=0.25$, $\beta=0.5$

```
#fig,ax = gradient_descent_plot(alpha,beta)
fig0,ax0 = gradient_descent_plot(0.25,0.5)
plt.savefig('fig1.png')
```




Below is the code to perform gradient descent with $\alpha=0.1$, $\beta=0.5$

```
fig1,ax1 = gradient_descent_plot(0.01,0.5)
plt.savefig('fig2.png')
```

Gradient Descent, $\alpha = 0.01$, $\beta = 0.50$

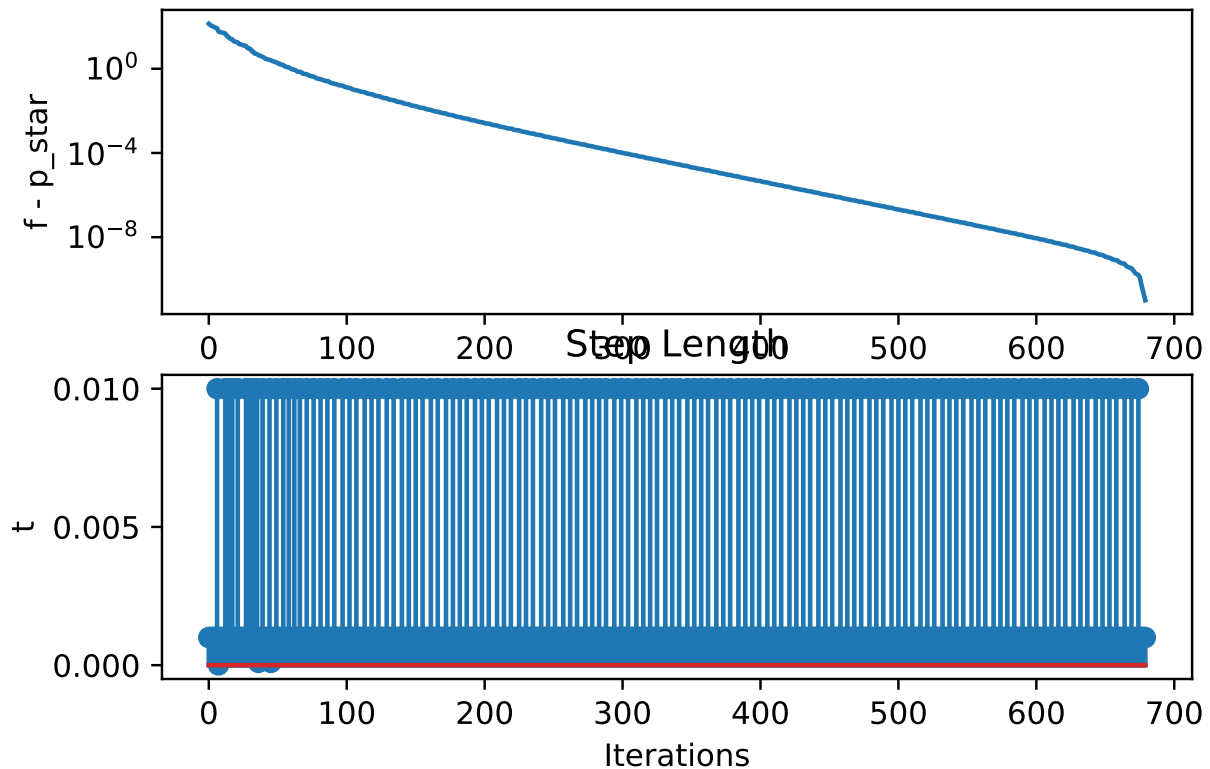
Convergence to Optimal Solution



The decrease in alpha eliminates some of the larger step size changes, and yields a smoother but slightly less accurate convergence in the bounds of the iterations. Let's try playing with beta a little bit now.

```
fig2, ax2 = gradient_descent_plot(0.01, 0.1)
plt.savefig('fig3.png')
```

Gradient Descent, $\alpha = 0.01$, $\beta = 0.10$ Convergence to Optimal Solution



Great! Our convergence has gotten even more accurate. Notice also the step size increase.

1.2 Part b - repeat using newton's method

```
def hess_f(x):
    s1 = 1/(1-np.dot(A,x))
    s2 = np.diag(np.power(s1,2)[: ,0])
    s3 = np.dot(s2,A)
    s4 = np.dot(np.transpose(A),s3)
    s5 = np.power(1+x,2)
    s6 = np.diag(1/s5)
    s7 = np.power(1-x,2)
    s8 = np.diag(1/s7)
    return s4+s6+s8

def newton_method(iterations,alpha,beta,eps):
    y = np.array([])
    s = np.array([])
    x = np.zeros(n)
    x = x.reshape(n,1) # Make sure x is a vector
    # Repeat
    # Step 1
    hf = hess_f(x)
    gf = grad_f(x)
    xnt = -np.linalg.solve(hf,gf)
    dec = -np.dot(np.transpose(gf),xnt)
    p = f(x)
    print(dec)
    for i in range(iterations):
```

```

        # Step 1
        y = np.append(y, f(x))

        # Step 2
        if dec/2 <= eps:
            break

        t = backtrack(x, xnt, alpha, beta)
        s = np.append(s, t)

        # Step 3
        x = x + t*xnt

        hf = hess_f(x)
        gf = grad_f(x)
        xnt = -np.linalg.solve(hf, gf)
        dec = -np.dot(np.transpose(gf), xnt)

        p = f(x)

    return y, s, p

def newton_method_plot(alpha, beta, eps):
    # Run gradient descent algorithm
    y, s, p = newton_method(i_max, alpha, beta, eps)

    fig, ax = plt.subplots(2, 1)

    ax[0].semilogy(y-p)
    ax[0].set_title('Convergence to Optimal Solution')
    ax[0].set_xlabel('Iterations')
    ax[0].set_ylabel('f - p_star')

    ax[1].stem(s)
    ax[1].set_title('Step Length')
    ax[1].set_xlabel('Iterations')
    ax[1].set_ylabel('t')

    fig.suptitle('Newton Method, alpha = %1.2f, beta = %1.2f'%(alpha, beta), fontsize=12)

    return fig, ax

```

1.2.1 Figures and analysis

Below is the code to perform newton method with alpha=0.25, beta=0.5

```

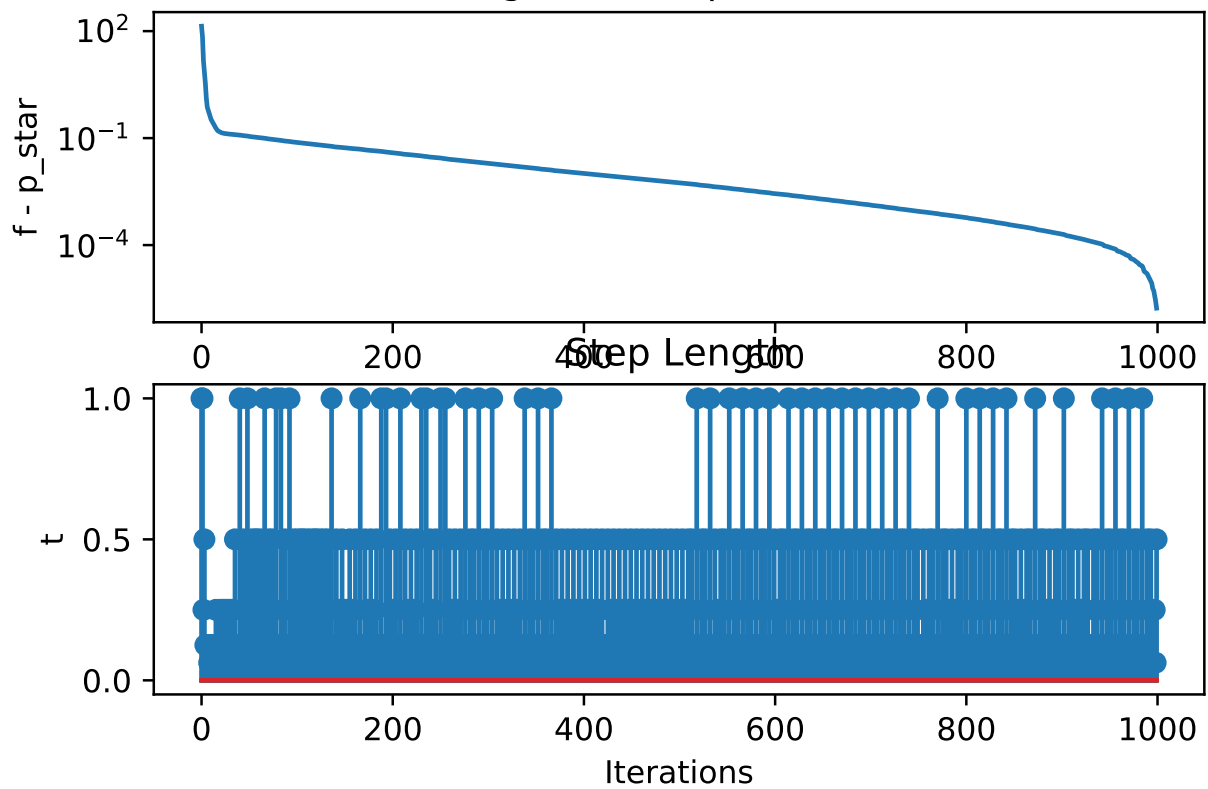
#fig, ax = newton_method_plot(alpha, beta)
fig3, ax3 = newton_method_plot(0.25, 0.5, eps)
plt.savefig('fig4.png')

```

```
[[101.82381699]]
```

Newton Method, $\alpha = 0.25$, $\beta = 0.50$

Convergence to Optimal Solution




Below is the code to perform newton with $\alpha=0.1$, $\beta=0.5$

```
fig4,ax4 = newton_method_plot(0.01,0.5,eps)
plt.savefig('fig5.png')
```

```
| [[101.82381699]]
```

Newton Method, $\alpha = 0.01$, $\beta = 0.50$

Convergence to Optimal Solution



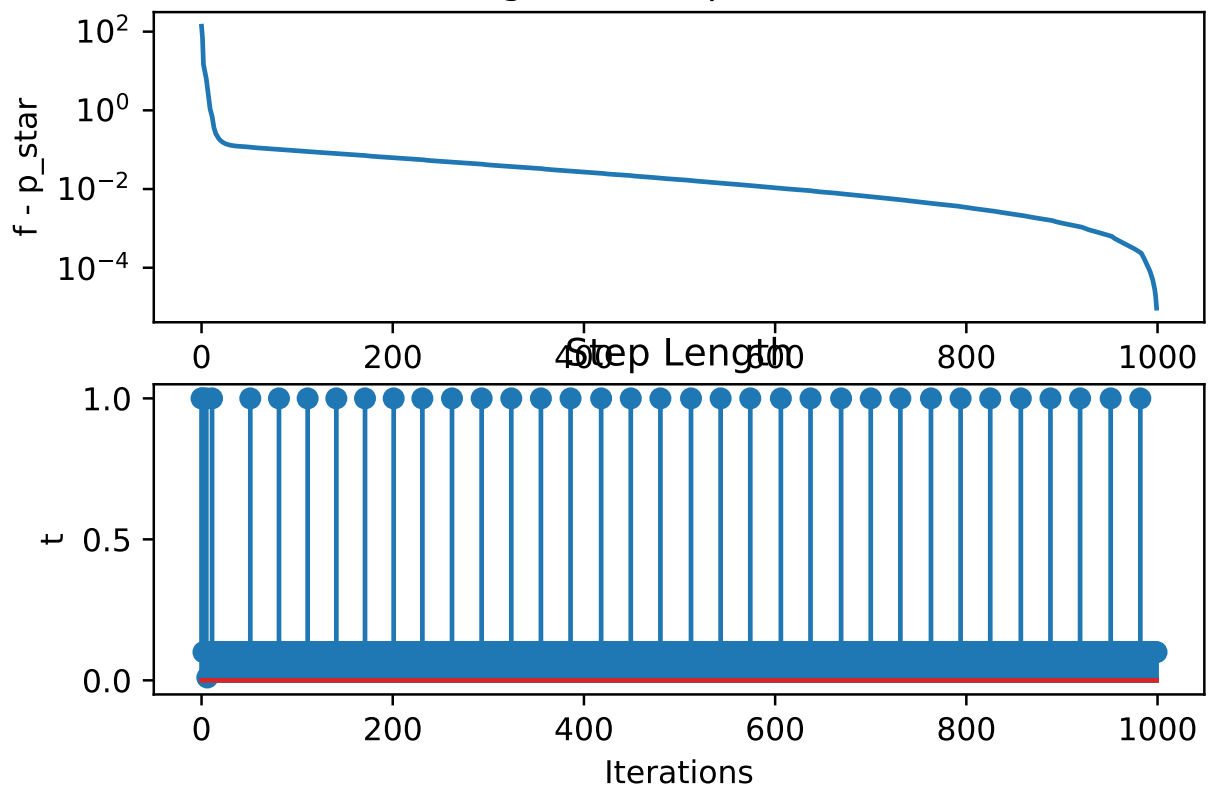
Let's try playing with beta a little bit now

```
fig5,ax5 = newton_method_plot(0.01,0.1,eps)
plt.savefig('fig6.png')
```

```
| [[101.82381699]]
```


Newton Method, $\alpha = 0.01$, $\beta = 0.10$

Convergence to Optimal Solution



Notice how in comparison to the gradient descent method, the newton method converges much more quickly than the gradient descent